



moe_unfrozen_8exp_sparse_b128 [FAILED]

$E=8.0 \cdot k=1.0 \cdot \text{batch}=128.0 \cdot \text{frozen=no} \cdot \text{FAILED} \cdot \text{role=Negative} - 8 \text{ experts}$


(a) Validation BER




(b) Noisy validation BER




(c) Expert max use



(d) Train-val load L1



(e) Validation expert loads at epoch 20



Best: 0.40% @ ep16 · Final: 1.36% @ ep20