

RQ4: MoE inference cost vs HiDDeN baseline (measured)
Sparse top-k=1: only one expert decoded per image regardless of expert count.

